

1 PROJECT 1: Regression of function data

The data stems from a non linear function. Hence, a non-linear transfer function is required. Bearing in mind the results of Lehno *et al.*, I find that the `tansig` (non-polynomial) function for the hidden layers must allow a function approximation if the number of neurons is sufficient. For the output layer, I use a `purelin` transfer function, which does not restrict the output values to a particular range of values. In most cases, a single layer neural network is sufficient to reproduce a function (in theory, it is sufficient for *all* functions if the number of neurons is large enough). Hence, we start with a single layer. Since the typical Euclidian distance between data points is small compared to the typical spacial structure dimension in addition to the fact that the data has not been jittered, the number of neurons can be taken to be relatively high without overfitting. To get an estimate of the number of neurons, we refer to Fig. 1. We note that by considering the validation set performance, the optimal number of neurons is between 25 – 40. Note that at this stage, we may not consider the test set performance. We opt for 35 neurons. Note that to generate Fig. 1, we did not use the validation set in a stopping criterion. This is because we want to *detect* overfitting, rather than *avoiding* it. The latter is the subject of the next paragraph.

First, we discuss the purpose of the data sets. Training is done on the training set, meaning that the network-parameters are varied such that it reproduces the training set with an overall error which is as small as possible. When the network starts to overfit the training data, the error on the validation set typically begins to rise. This indicates that the network is adjusted in such a way that it minimizes the error with respect to the training set, while having bad generalization properties in the regions where no training data is available. Therefore, when the validation error increases for a specified number of iterations (`net.trainParam.max_fail` = 20), the training is stopped, and the weights and biases at the minimum of the validation error are returned. The test set only serves as a final check and cannot be used in the training and validation process.




Figure 1: Left: surface and data points in the training set. Right: Performance on the training, validation and test set as a function of the number of neurons in the hidden layer.

The MSE on the test set of the resulting NN is 5.5×10^{-5} . We observe the pattern that the network has slightly higher values when the test set interpolation is large, and smaller values when the interpolations function is small. Hence, more extremal values are obtained.

To improve the results, we can use all of the provided data, instead of a subset of 3000. Also, the data division does not need to be equal for all sets. As a rule of thumb,




Figure 2: Performance on the training, validation and test set as a function of the number of epochs in the training process.

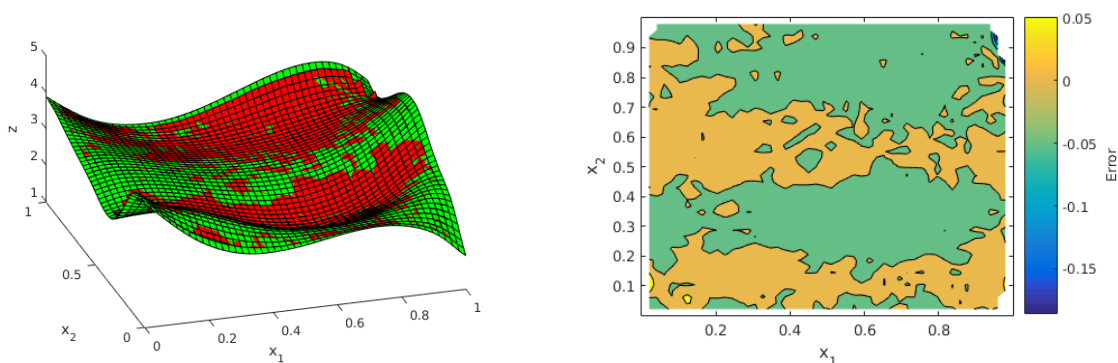


Figure 3: Left: Neural network (green) and test data (red) surface. The surfaces are almost indistinguishable. Right: Contour plot of the test set error.

one generally assumes 70:15:15 (or 60:20:20) for the train:validation:test set ratios. This changes the MSE test error to 4.1×10^{-8} .

2 PROJECT 2: Classification of wine data

Since it is not given whether or not the wine data is linearly separable, we use a non-linear model. We now use a `tansig` transfer function in the output layer, since this restricts the output to the domain $[-1, 1]$ and the target output is in the format of ± 1 . We rescale the attributes to standardized values in order to make it independent of the scale of the data attributes. 15 neurons in the hidden layer is optimal to classify the validation data. Using this architecture, a CCR of 0.60 and 0.72 is obtained for the validation and test set respectively. This might indicate that the two wine classes have a significant overlap in their attribute space.

Next, the data is projected onto its lower dimensional principle component basis and reconstructed afterwards. The eigenvalue spectrum of the covariance matrix is depicted in Fig. 4. Prior to computing the covariance matrix, we again rescale the data such that all attributes have zero mean and unit standard deviation. The eigenvalue spectrum does not show a clear dominant principle component. Hence, there is no clear indication that the data lives on a low-dimensional subspace. We take a reconstruction error of 10%, which corresponds with an 8 dimensional basis. 11 hidden neurons minimize the CCR of the validation set, which is less than before, as expected. Using this approach, a validation




Figure 4: Eigenvalues of the 11×11 covariance matrix. The red line indicates the cumulative sum of remaining $(k+1):end$ eigenvalues, which is proportional to the reconstruction error.

and test set CCR of 0.75 and 0.72 is obtained. Hence, there is a significant CCR increase for the validation set, while test set CCR is almost constant.

Hence, by projecting the data onto a lower dimensional space, a significant classification improvement of approximately 14% of the validation set is realized. However, the result is not true in general, since such a large increase is not found for the test set. This is related to the fact that the principle component basis is that of the training set, and is not updated if new data is included.

We remark that the possibility of overfitting is very real in this classification assignment.

3 PROJECT 3: Character recognition

3.1 Hopfield network




Figure 5: First 10 of the 32 letters that are used in the character recognition exercise.

In this section we train a Hopfield network to reconstruct the states in Fig. 5 from a distorted version of these characters. Hopfield networks are generated by providing it the undistorted images.




Figure 6: Incorrectly reconstructed characters (top row) with corresponding input (bottom row).

Figure 6 illustrates a number of incorrectly reconstructed states with corresponding undistorted input. The ‘e’ and ‘a’ characters are often interchanged in the reconstruction. This is related to the fact that these characters have many pixels with the same value, and hence a similar shape. This means that the ridge in the energy function between these two attractor states is lower than the 3-pixel distortion. The incorrectly returned ‘j’ character shows a spurious state which is a linear combination of multiple characters. These cases are obtained when allowing for a sufficiently large number of time steps (1000). This

allows one to observe the spurious states that are inherent to the system, rather than states that have not yet converged.




Figure 7: Total pixel error (red), normalized over the number of states generated and total pixels in an image (35), and the theoretical prediction curve based on the Hebb rule (blue).

We now determine how number of patterns that are stored influence the number of erroneously restored characters. For each case of P stored patterns, we create a Hopfield network and generate a sufficient number of 3-pixel distorted images. These are then reconstructed, after which we calculate the number of pixels that do not correspond to the original image. The results are shown in Fig. 7. One can observe a steep rise in the reconstructed pixel error after $P = 17$. Hence, 17 stored patterns is the critical loading capacity of the network. Naturally, when the number of distorted pixels is increased, this number is expected to go down. For P larger than the critical loading capacity, the large number of spurious states and the small basins of attraction of attractor state do not allow for a reliable character restoration.

One can theoretically predict the dependence of the reconstruction error as a function of the number of patterns stored, using Hebb's rule for uncorrelated patterns. The prediction curve is also shown in Fig. 7. If we assume a large P and large N , we obtain an estimated P_{error}

$$P_{\text{error}} = \frac{1}{\sqrt{2\pi}\sigma} \int_1^{+\infty} e^{-\frac{x^2}{2\sigma^2}} dx, \quad (1)$$

where $\sigma = \sqrt{P/N}$. Hence, a first approximation is that P and N are large, which is not actually the case here. However, both the prediction and simulation show the same (expected) behaviour: first a flat P dependence, followed by a steep rise. One can estimate the storage capacity from

$$P_{\text{max}} = N/(4 \log N) = 2.46. \quad (2)$$

Hence, for 2 or less stored patterns, the Hopfield should be able to perfectly reconstruct the distorted images.

One way to resolve the issue of incorrectly reconstructed images, is to increase the number of pixels in each image. In the case at hand, each image is determined by 35 pixels. Hence, it was expected that the critical loading capacity would be relatively low. By increasing the number of pixels in an image, each character has more attributes by which it is characterized. This is the obvious alternative to the current Hopfield network. In the next subsection we discuss another alternative.

3.2 Alternative solution to character recognition

The Hopfield network for character recognition is “only implemented in MATLAB for historical reasons”. Nowadays, many more, and better ANNs are available on the market. Not only can one choose the architecture of the network, but also the error determination can be chosen to optimize the problem. Another choice is which input is used to train the network. It has been shown that (as expected), a network performs better in character restoration if it is trained with distorted input images. This is clearly illustrated in an example by Matlab ¹.

The opted network architecture is as follows: we generate a feedforward neural network with one hidden layer of 25 neurons and an output layer of dimension equal to the number of patterns P that are to be stored. In the hidden layer, we use a `tansig` transfer function, while in the output layer, we use the `softmax` function. The final output is such that we put all output components to 0, except for the one with the maximum value, which is 1. Hence, we use one-hot encoding of the stored characters. This means that every character out of the P stored characters is determined by a P -dimensional vector of zeros, except for a one in a unique place. Hence, also the actual characters must be stored such that one can return a character instead of this vector. Note that the input still takes the 35 pixels as in the original Hopfield network.

One-hot encoding is feasible if the number of patterns stored is not too large, since the size of the output layer is equal to the size of the alphabet to be stored.

In the previous subsection we showed that the “e” and “a” are difficult for the Hopfield network to restore correctly, since these characters have very similar features. Hence, if one would create a network that uses an $N = 35$ dimensional output, similar problems would be faced, since these characters are only separated by a small distance in output space. In one hot-encoding, the distance between two patterns in output space is $\sqrt{2}$, independent of whether they have similar patterns.

We train our network with 100 distorted images of each letter to be stored. We vary the number of patterns and determine the number of incorrectly restored patterns as a function of the number of patterns stored. We use a constant 25 hidden neurons for practical reasons. The results are shown in Fig. 8. A steady rise is observed as expected.

Naturally, the networks’ performance can easily be improved by increasing the number of neurons and/or hidden layers. In addition, approximately 6500 possible 3-pixel distortions of each letter. Hence, in theory all possibilities can be generated to form the training batch. However, substantial computing resources are required if batch-training is required. We will not go into further detail, since these calculations are hard to perform on a simple laptop.

¹<http://nl.mathworks.com/help/nnet/examples/character-recognition.html>




Figure 8: Percentage of false reconstructed images (not pixels as before) as a function of the number of patterns stored.

A Matlab code that implements the solutions

A.1 Regression (see Section 0.1)

A.1.1 Main code

```

1 clear all; close all; clc;
2
3 %% LOAD THE DATA FROM TOLEDO
4
5 % load the data from Toledo
6 load 'data/Data_Problem1_regression.mat'
7
8 % generate my data from studentnr 0639870
9 d1 = 9;
10 d2 = 8;
11 d3 = 7;
12 d4 = 6;
13 d5 = 3;
14 T = (d1*T1 + d2*T2 + d3*T3 + d4*T4 + d5*T5)/(d1 + d2 + d3 + d4 + d5);
15 X = [X1 X2]; % useful for running sim
16 N = size(T,1);
17
18 % a check
19 assert((size(X1,1) == N) && (size(X2,1) == N));
20
21 %% DIVIDE THE DATA IN SUBSETS
22
23 % randomize the data, shuffle it, also transpose it
24 shuffledInd = randperm(size(T,1));
25 X1 = X1(shuffledInd);
26 X2 = X2(shuffledInd);
27 T = T(shuffledInd);
28
29 % divide the data in subsets (use the first 3000, they are already

```

```

    shuffled)
30 rng(1);
31 useN = 3000;
32 [trainInd, valInd, testInd] = dividerand(useN, 1/3, 1/3, 1/3);
33
34 % easier variables for training the network
35 xtrain = [X1(trainInd)'; X2(trainInd)']; % input
36 ytrain = T(trainInd, :)'; % target
37 xval = [X1(valInd)'; X2(valInd)'];
38 yval = T(valInd)';
39 xtest = [X1(testInd)'; X2(testInd)'];
40 ytest = T(testInd)';
41
42 %% VIZUALISE THE DATA
43 % create a regular grid for the interpolant (which is not the net)
44 ndim = 50;
45 [Xmesh, Ymesh, ~, ~] = meshinterpolate(X1, X2, T, ndim);
46 % now interpolate
47 trainInterpolate = TriScatteredInterp(X1(trainInd), X2(trainInd), T(
    trainInd)); % plot the interpolator
48 Zmesh = trainInterpolate(Xmesh, Ymesh);
49 % show the surface
50 figure;
51 surface(Xmesh, Ymesh, Zmesh); hold on;
52 % also show the data
53 scatter3(X1(trainInd), X2(trainInd), T(trainInd), '.', '
    MarkerEdgeColor', 'red', 'MarkerFaceColor', 'red');
54 xlabel('x_1'); ylabel('x_2'); zlabel('z');
55 colorbar; h = colorbar; ylabel(h, 'z');
56 savefig('train_surface.fig'); hold off;
57
58 % show that the data is well distributed and divided randomly
59 figure;
60 surface(Xmesh, Ymesh, Zmesh); hold on; % plot the interpolator
61 scatter3(X1(trainInd), X2(trainInd), T(trainInd), 15, '
    MarkerFaceColor', 'g'); hold on;
62 scatter3(X1(valInd), X2(valInd), T(valInd), 15, 'MarkerFaceColor', [1
    .5 0]); hold on;
63 scatter3(X1(testInd), X2(testInd), T(testInd), 15, 'MarkerFaceColor',
    'r'); hold off;
64
65 %% CREATE THE VALIDATION PLOT
66
67 % create some vectors to plot
68 nhvals = [1, 5, 10, 15, 20, 25, 30, 35, 40, 45, 50];
69 mseVal = zeros(length(nhvals), 1);
70 mseTrain = zeros(length(nhvals), 1);
71 mseTest = zeros(length(nhvals), 1);
72
73 for nhIt = 1:length(nhvals)
74

```



```

75     % change the number of neurons in the hidden layer
76     nh = nhvals(nhIt);
77     disp(['Training with ' num2str(nh) ' neurons in hidden layer']);
78
79     rng(1); % same seed for all
80
81     % train the network
82     net = feedforwardnet(nh, 'trainlm');
83     net.divideFcn = 'dividetrain'; % Use the whole training set for
        training
84     net.layers{1}.transferFcn = 'tansig';
85     net.layers{2}.transferFcn = 'purelin';
86     net.trainParam.showWindow=0;
87     net.trainParam.epochs = 5000;
88
89     % train the network
90     [net, tr] = train(net, xtrain, ytrain, 'UseParallel', 'yes');
91     nntraintool('close');
92
93     % show number of epochs used
94     disp(['Used ' num2str(tr.num_epochs) ' epochs']);
95
96     % make predictions for the data sets
97     ytrainPred = sim(net, xtrain);
98     yvalPred = sim(net, xval);
99     ytestPred = sim(net, xtest); % plot this for completeness
100
101     % calculate the mse for these points
102     perfTrain = perform(net, ytrain, ytrainPred);
103     perfVal = perform(net, yval, yvalPred);
104     perfTest = perform(net, ytest, ytestPred);
105
106     % store it
107     mseTrain(nhIt) = perfTrain;
108     mseVal(nhIt) = perfVal;
109     mseTest(nhIt) = perfTest;
110 end
111
112 %% SELECT A NETWORK ARCHITECTURE FROM THE RESULTS
113 nhfinal = 35;
114
115 %% SHOW PERFORMANCE ON VALIDATION SET
116 figure
117 semilogy(nhvals, mseTrain, 'Color', 'b', 'LineWidth', 2); hold on;
118 semilogy(nhvals, mseVal, 'Color', 'g', 'LineWidth', 2); hold on;
119 semilogy(nhvals, mseTest, 'Color', 'r', 'LineWidth', 2);
120 xlabel('Neurons in the hidden layer'); ylabel('MSE');
121 savefig('performance_val.fig');
122
123 %% PLOT THE RESULT OF THE OPTED ARCHITECTURE
124 % train the network

```

```

125 rng(1);
126 net = feedforwardnet(nhfinal,'trainlm');
127 net.divideFcn = 'divideind'; % also give it access to the other sets
    for plotting
128 net.divideParam.trainInd = trainInd;
129 net.divideParam.valInd = valInd;
130 net.divideParam.testInd = testInd;
131 net.layers{1}.transferFcn = 'tansig'; % hidden layer
132 net.layers{2}.transferFcn = 'purelin'; % output layer
133 net.trainParam.epochs = 5000; % set really high, so it can decide
    itself
134 net.trainParam.max_fail = 50; % set really high, so it can decide
    itself
135 [net,tr] = train(net, xtrain, ytrain);
136
137 % calculate its output on the meshgrid
138 ZmeshNN = net([Xmesh(:) Ymesh(:)]');
139 ZmeshNN = reshape(ZmeshNN, size(Xmesh,1), size(Xmesh,2));
140
141 % map the interpolator of the test set
142 testInterpolant = TriScatteredInterp(X1(testInd), X2(testInd), T(
    testInd));
143 ZmeshTest = testInterpolant(Xmesh, Ymesh);
144
145 % plot it together with the interpolant of the TEST set
146 figure
147 surface(Xmesh, Ymesh, ZmeshNN, 'FaceColor', 'g'); hold on; % plot the
    NN
148 surface(Xmesh, Ymesh, ZmeshTest, 'FaceColor', 'r'); % plot the
    interpolator
149 xlabel('x_1'); ylabel('x_2'); zlabel('z');
150 savefig('NN_and_testsurf.fig');
151
152 % make the error plot
153 ZmeshError = ZmeshNN - ZmeshTest;
154 figure
155 contourf(Xmesh, Ymesh, ZmeshError);
156 xlabel('x_1'); ylabel('x_2'); colorbar;
157 h = colorbar; ylabel(h, 'Error');
158 savefig('NN_test_error.fig');
159
160 % compute MSE on the test set
161 mseTest = sum((net(xtest)-ytest).^2)/size(ytest,2);
162 fprintf('MSE on the test set: %.2E\n', mseTest);
163
164 % plot the regression for the test set
165 figure;
166 plotregression(sim(net, xtest), ytest);
167
168
169

```

```

170 %% USE ALL THE DATA TO SEE WHAT HAPPENS, ALSO USE BETTER RATIOS
171 rng(1);
172 useN = size(T,1);
173 [trainInd, valInd, testInd] = dividerand(useN);
174 % easier variables for training the network
175 xtrain = [X1(trainInd)'; X2(trainInd)']; % input
176 ytrain = T(trainInd,:)'; % target
177 xval = [X1(valInd)'; X2(valInd)'];
178 yval = T(valInd)';
179 xtest = [X1(testInd)'; X2(testInd)'];
180 ytest = T(testInd)';
181 % train the network
182 rng(1);
183 net = feedforwardnet(nhfinal, 'trainlm');
184 net.divideFcn = 'divideind'; % also give it access to the other sets
    for plotting
185 net.divideParam.trainInd = trainInd;
186 net.divideParam.valInd = valInd;
187 net.divideParam.testInd = testInd;
188 net.layers{1}.transferFcn = 'tansig'; % hidden layer
189 net.layers{2}.transferFcn = 'purelin'; % output layer
190 net.trainParam.epochs = 5000; % set really high, so it can decide
    itself
191 net.trainParam.max_fail = 50; % set really high, so it can decide
    itself
192 [net, tr] = train(net, xtrain, ytrain, 'UseParallel', 'yes');
193 % calculate its output on the meshgrid
194 ZmeshNN = net([Xmesh(:) Ymesh(:)]');
195 ZmeshNN = reshape(ZmeshNN, size(Xmesh,1), size(Xmesh,2));
196 % map the interpolator of the test set
197 testInterpolant = TriScatteredInterp(X1(testInd), X2(testInd), T(
    testInd));
198 ZmeshTest = testInterpolant(Xmesh, Ymesh);
199 % plot it together with the interpolant of the TEST set
200 figure
201 surface(Xmesh, Ymesh, ZmeshNN, 'FaceColor', 'g'); hold on; % plot the
    NN
202 surface(Xmesh, Ymesh, ZmeshTest, 'FaceColor', 'r'); % plot the
    interpolator
203 xlabel('x_1'); ylabel('x_2'); zlabel('z');
204 % make the error plot
205 ZmeshError = ZmeshNN - ZmeshTest;
206 figure
207 contourf(Xmesh, Ymesh, ZmeshError);
208 xlabel('x_1'); ylabel('x_2'); colorbar;
209 h = colorbar; ylabel(h, 'Error');
210 % compute MSE on the test set
211 mseTest = sum((net(xtest)-ytest).^2)/size(ytest,2);
212 fprintf('MSE on the test set: %.2E\n', mseTest);

```

A.2 Classification (see Section 0.2)

A.2.1 Main code

```

1 clear all; close all; clc;
2
3 %% LOAD THE DATA FROM TOLEDO
4
5 % generate my data from studentnr 0639870
6 % digit 0 gives me (C+,C-) = (5,6) of the white wine
7 % these classes are in the last column
8 datatable = importdata('data/winequality-white.csv');
9 data = datatable.data;
10 pos = data(data(:,end) == 5,:); % 5 is positive
11 neg = data(data(:,end) == 6,:); % 6 is negative
12 Npos = size(pos,1);
13 Nneg = size(neg,1);
14
15 % put these in more useful training formats (all in one)
16 % input are all columns, except for the last one, which is target
17 % we set the target values to +/-1 instead of 5 and 6
18 X = [pos(:,1:(end-1)) ; neg(:,1:(end-1))]';
19 T = [ones(Npos,1) ; -ones(Nneg,1)]';
20 N = Npos + Nneg;
21
22 % shuffle and divide the data
23 rng(1);
24 [trainInd, valInd, testInd] = dividerand(N); % default is 0.7, 0.15,
    0.15
25 stdX = mapstd(X);
26
27
28 %% CREATING A NEURAL NETWORK TO CLASSIFY THE DATA
29 % create a network and train it
30 rng(1);
31 net = feedforwardnet(15, 'trainlm');
32 % classification values are between -1 and 1, hence, we can use the
    tangent
33 % sigmoid function in the output layer as well
34 net.layers{1}.transferFcn = 'tansig'; % hidden layer
35 net.layers{2}.transferFcn = 'tansig'; % output layer
36 net.divideFcn = 'divideind';
37 net.divideParam.trainInd = trainInd;
38 net.divideParam.valInd = valInd;
39 net.divideParam.testInd = testInd;
40 net.trainParam.max_fail = 50; % may vary this
41 net.trainParam.min_grad = 10^-15; % may vary this
42 net = train(net, stdX, T);
43
44 %% PERFORMANCE CHECKS
45 predVal = sim(net, X(:, valInd));
46 CCRval = sum(sign(predVal) == T(valInd))*100/length(valInd);

```

```

47 predTest = sim(net, stdX(:, testInd));
48 CCRtest = sum(sign(predTest) == T(testInd))*100/length(testInd);
49 fprintf('CCRval= %f and CCRtest %f .\n', CCRval, CCRtest);
50
51 %% PCA
52 % preprocess the data to get zero mean = 0 and stddev = 1 for all
53 % properties
54 Ntrain = size(X(:, trainInd), 2); % number of training data points
55 [~, ~, eigvals, ~, ~, ~] = doPCA(X(:, trainInd)', 11); % use all 11
56
57 % plot the eigenvalues
58 figure;
59 bar(eigvals/max(eigvals)); hold on; % shows that one needs 'only' 10
    basis vectors
60 plot(1:11, 1-cumsum(eigvals/sum(eigvals)), 'r-', 'LineWidth', 2)
61 ylabel('\lambda_k'); xlabel('k');
62 axis([0 12 0 1]);
63 savefig('eigenvalues.fig');
64
65 % we project the vectors onto the restricted eigenbasis (columns of
    eigvecs)
66 numBasisVecs=8; % choose the number of eigenvectors
67 [PCABasis, redXTrain, eigvals, meanTrain, stddevTrain, stdXTrain] = doPCA(
    X(:, trainInd)', numBasisVecs);
68
69 % project also validation and test set, but first standardize them,
    use
70 % part of my doPCA function for this
71 % note that we project with the PCA basis of the training set, as
    required
72 % by the assignment
73 [~, ~, ~, meanVal, stddevVal, stdXVal] = doPCA(X(:, valInd)', numBasisVecs);
74 [~, ~, ~, meanTest, stddevTest, stdXTest] = doPCA(X(:, testInd)',
    numBasisVecs);
75 redXVal = stdXVal*PCABasis;
76 redXTest = stdXTest*PCABasis;
77
78 % now reconstruct
79 redXTrain = redXTrain*PCABasis';
80 redXVal = redXVal*PCABasis';
81 redXTest = redXTest*PCABasis';
82
83 % create a new input matrix X from these components
84 %X = zeros(numBasisVecs, N);
85 X(:, trainInd) = redXTrain';
86 X(:, valInd) = redXVal';
87 X(:, testInd) = redXTest';
88
89 % we can now use them for training
90 %% CREATING A NEURAL NETWORK TO CLASSIFY THE DATA
91 % create a network and train it

```

```

92 rng(1);
93 net = feedforwardnet(11, 'trainlm');
94 net.layers{1}.transferFcn = 'tansig'; % hidden layer
95 net.layers{2}.transferFcn = 'tansig'; % output layer
96 net.divideFcn = 'divideind';
97 net.divideParam.trainInd = trainInd;
98 net.divideParam.valInd = valInd;
99 net.divideParam.testInd = testInd;
100 net.trainParam.max_fail = 50; % may vary this
101 net.trainParam.min_grad = 10^-15; % may vary this
102 [net, tr] = train(net,X,T);
103
104 %% PERFORMANCE CHECKS
105 predVal = sim(net, X(:,valInd));
106 CCRval = sum(sign(predVal) == T(valInd))*100/length(valInd);
107 predTest = sim(net, X(:,testInd));
108 CCRtest = sum(sign(predTest) == T(testInd))*100/length(testInd);
109 fprintf('CCRval= %f and CCRtest %f .\n',CCRval,CCRtest);

```

A.2.2 PCA related

```

1 function [E,z,d,meanVec,stddevVec,stdX] = doPCA(x, q)
2 % DOPCA Do a Principle Component Analysis on a data set x
3 % [E,z,mean,stddev] = DOPCA(x,q) with x a dataset with each row a
  data entry ,
4 % performs PCA with result z and matrix E.
5 % this means that z is x in a reduced basis
6 % q is the reduced dimension (dimension of z)
7 % d is the eigenvalue vector
8 % meanVec stddevVec contain the mean and standard deviation vector
  of the
9 % original data
10 % stdX contains the standardized set in unreduced space
11
12 % get the dimension p of x and the number of datapoints N
13 p = size(x,2);
14 assert( q <= p );
15 N = size(x,1);
16
17 % calculate the mean for all p data properties
18 meanVec = mean(x);
19 stddevVec = std(x);
20 %stddevVec = ones(1,p);
21
22 % rescale and shift with these vectors
23 stdX = zeros(N, 11);
24 for i = 1:N
25     stdX(i,:) = x(i,:) - meanVec; % shift
26     stdX(i,:) = stdX(i,:) ./ stddevVec; % rescale
27 end
28
29 % calculate the covariance matrix of the standardized data set

```

```

30 V = cov(stdX);
31
32 % calculate the eigenvectors and eigenvalues
33 % E is the matrix with columns the eigenvectors
34 [E,d] = eig(V);
35 d = diag(d);
36
37 % now sort them in descending order and take only q of them
38 [d, indices] = sort(d, 'descend');
39 d = d(1:q);
40 E = E(:,indices(1:q));
41
42 % reduce the data set my multiplying with this matrix
43 z=stdX*E;

```

A.3 Classification (see Section 0.2)

A.3.1 Main code

```

1 close all; clc; clear all;
2
3 %% CREATE THE ATTRACTOR STATES
4 % get all the letters of the alphabet in CAPITALS
5 [ALPHABET, ~]=prprob();
6
7 % now the unique lower case letters of my name
8 name=GenerateName;
9
10 % add the ALPHABET and name
11 allLetters=[name';ALPHABET']';
12
13 % rescale from -1 to 1 instead of 0 and 1
14 allLetters = 2*allLetters - 1;
15
16 % plot the first 10 letters of my data set
17 figure;
18 colormap(gray);
19 for letterNr=1:10
20     subplot(2,5,letterNr);
21     imagesc(reshape(allLetters(:,letterNr),5,7)', 'CDataMapping', '
        scaled'); % 5x7 bit maps transposed to a 7x5
22     set(gca, 'xtick', []); set(gca, 'xticklabel', []); set(gca, 'ytick',
        [], []); set(gca, 'yticklabel', [], []);
23 end
24 savefig('letters.fig');
25 hold off; close all;
26
27 %% CREATE A HOPFIELD RECURRENT NETWORK, TYPE 1
28 % type 1 retrieves 5 first letters
29 T = allLetters(:,1:5);
30 net = newhop(T);
31

```

```

32 %% DISTORT AND RETRIEVE IMAGES
33 % check the correct retrieval rate, output states that are spurious
34 % even though we can do the following exactly, we do a
35 % simulation. We create 1000 distorted images of each of the 5
    letters and
36 % use the Hopfield network to retrieve the original states. If the
    number
37 % of attempts is large enough, we should find all spurious states.
    Note
38 % however that there are  $35!/(3!32!)=6545$  possible distorted images
    of
39 % each of the letters
40 % Compared to the assignment on Hopfields, the distorted image now
    has
41 % discrete values for the pixels. Hence, it's now more feasible to
    end up
42 % in a spurious state
43
44 Nwrong = 0;
45 timesteps = 1000;
46 for letterNr = 1:5
47     fprintf('Start with letter nr : %i\n', letterNr)
48     letter = T(:, letterNr);
49     for it = 1:1000
50         distImage = DistortImage(letter);
51         [Y,~,~] = net({1 timesteps},{},{distImage});
52         if ~isequal(Y{end}, letter)
53             if (Nwrong ~= 0 && sum(ismember(Y{end}', wrongStates', '
                rows')==1) % avoid doubles
54                 fprintf('Same state.\n')
55             else
56                 Nwrong = Nwrong + 1;
57                 fprintf(' Wrong state at iteration: %i\n', it);
58                 wrongStates(:, Nwrong) = Y{end}; % add the state
59                 originalStates(:, Nwrong) = letter;
60             end
61         end
62     end
63 end
64
65 figure;
66 colormap(gray)
67 for wrongNr = 1:Nwrong
68     subplot(2, Nwrong, wrongNr);
69     imagesc(reshape(wrongStates(:, wrongNr), 5, 7)', 'CDataMapping', '
        scaled'); hold on;
70     subplot(2, Nwrong, Nwrong+wrongNr);
71     imagesc(reshape(originalStates(:, wrongNr), 5, 7)', 'CDataMapping', '
        scaled'); hold on;
72 end
73 hold all;

```



```

74 savefig('wrong_states.fig'); hold off;
75
76 %% MAPPING ERROR IFO P
77 Nit = 100;
78 % store number of wrong results
79 Nwrong = zeros(1,size(allLetters,2));
80 % loop over the number of stored patterns P
81 for P = 1:size(allLetters,2)
82     fprintf('Simulating with P = %i patterns stored.\n',P);
83     % take P attractors
84     T = allLetters(:,1:P);
85     % create a hopfield net
86     net = newhop(T);
87     % loop over Nit distorted images per letter and calculate the
        error
88     for letterNr = 1:P
89         letter = T(:,P);
90         for it = 1:Nit
91             distImage = DistortImage(letter);
92             % use the net to retrieve
93             [Y,~,~] = net({1 timesteps},{},{distImage});
94
95             % clip
96             Y{end} = sign(Y{end});
97             % check if it's the correct one
98             Nwrong(P) = Nwrong(P) + sum(abs(Y{end} - letter));
99         end
100     end
101     Nwrong(P) = Nwrong(P) / (Nit*P*size(allLetters,1)); % normalize
        over number of states that we generated and nr of pixels
102 end
103
104 % estimate also with Hebb rule
105 figure;
106 sigmas = sqrt((1:size(allLetters,2))/size(allLetters,1));
107 mus = zeros(1,size(allLetters,2));
108 Perr = ones(1,size(allLetters,2))-normcdf(ones(1,size(allLetters,2)),
        mus,sigmas);
109 %plot it
110 plot(1:size(allLetters,2),Nwrong,'r-','LineWidth',2); hold on;
111 plot(1:size(allLetters,2),Perr,'b-','LineWidth',2);
112 legend('Pixel error','P_{error}')
113 ylabel('Error');
114 xlabel('Number of patterns stored');
115 savefig('Error_ifo_P.fig');

```

A.3.2 Image distortion

```

1 function distImage = DistortImage(image)
2
3 % take three random numbers between 1 and 35
4 indices = randsample(35,3); % unique, no replacement

```

```

5
6 % now get those pixels and switch them
7 distImage = image;
8 distImage(indices) = -distImage(indices);

```

A.3.3 Name generation

```

1 function lowercasename = GenerateName()
2
3
4 j = [...
5     0 0 0 0 0 ...
6     0 0 0 0 1 ...
7     0 0 0 0 0 ...
8     0 0 0 0 1 ...
9     0 0 0 0 1 ...
10    1 0 0 0 1 ...
11    0 1 1 1 0 ]';
12
13 a = [...
14     0 0 0 0 0 ...
15     0 0 0 0 0 ...
16     0 1 1 1 0 ...
17     0 0 0 0 1 ...
18     0 1 1 1 1 ...
19     1 0 0 0 1 ...
20     0 1 1 1 0 ]';
21
22 n = [...
23     0 0 0 0 0 ...
24     0 0 0 0 0 ...
25     1 0 1 1 0 ...
26     1 1 0 0 1 ...
27     1 0 0 0 1 ...
28     1 0 0 0 1 ...
29     1 0 0 0 1 ]';
30
31 e = [...
32     0 0 0 0 0 ...
33     0 0 0 0 0 ...
34     0 1 1 1 0 ...
35     1 0 0 0 1 ...
36     1 1 1 1 1 ...
37     1 0 0 0 0 ...
38     0 1 1 1 0 ]';
39
40 s = [...
41     0 0 0 0 0 ...
42     0 0 0 0 0 ...
43     0 1 1 1 1 ...
44     1 0 0 0 0 ...
45     0 1 1 1 0 ...

```

```

46     0 0 0 0 1 ...
47     1 1 1 1 0 ]';
48
49 y = [ ...
50     0 0 0 0 0 ...
51     0 0 0 0 0 ...
52     1 0 0 0 1 ...
53     1 0 0 0 1 ...
54     0 1 1 1 1 ...
55     0 0 0 0 1 ...
56     1 1 1 1 0 ]';
57
58 lowercasename = [j,a,n,e,s,y];

```

A.3.4 Alternative to Hopfield

```

1  close all; clc; clear all;
2
3  %% CREATE THE ATTRACTOR STATES
4  % get all the letters of the alphabet in CAPITALS
5  [ALPHABET, ~]=prprob();
6  % now the unique lower case letters of my name
7  name=GenerateName;
8  % add the ALPHABET and name
9  allLetters=[name';ALPHABET']';
10 % rescale from -1 to 1 instead of 0 and 1
11 allLetters = 2*allLetters - 1;
12
13 % loop over number of patterns stored
14 MaxNpatt=25;
15 icf = zeros(1,MaxNpatt);
16 %storage for some wrong letters, take one correct and one incorrect
   per
17 for P=1:size(icf,2)
18     %% CREATE THE FULL SET OF POSSIBILITIES
19     % make the set of the first 25 letters
20     letters = allLetters(:,1:P);
21
22     % make some possibilities of distortion of 3 pixels
23     % this is used for training
24     numDist = 100; % number of distorted images per letter
25     X = zeros(size(letters,1),size(letters,2)*numDist);
26     T = zeros(size(letters,2),size(letters,2)*numDist);
27     for il = 1:size(letters,2)
28         letter = letters(:,il);
29         for id = 1:numDist
30             X(:,(il-1)*numDist+id) = DistortImage(letter);
31             T(il,(il-1)*numDist+id) = 1; % only that one is 1
32         end
33     end
34
35     %% TRAIN A NEURAL NETWORK

```

```

36 % create network with 25 neurons
37 net = feedforwardnet(25);
38 net.layers{1}.transferFcn = 'tansig';
39 net.layers{2}.transferFcn = 'softmax';
40 net.trainParam.showWindow=0;
41 net = train(net, X, T, 'useParallel', 'yes'); % 25 dimensional
    output
42 % all letters have same distance of sqrt2
43
44 %% NOW CHECK ITS CAPABILITIES
45 Nit = 1000;
46 % store number of wrong results
47 Nwrong = zeros(1, size(letters, 2));
48
49 % loop over Nit distorted images per letter and calculate the
    error
50 parfor letterNr = 1:size(letters, 2)
51     letter = letters(:, letterNr);
52     fprintf('Starting with letter %i out of %i...\n', letterNr,
        size(letters, 2))
53     for it = 1:Nit
54         distImage = DistortImage(letter);
55         % use the net to retrieve
56         [Y, ~, ~] = net(distImage);
57
58         % one-hot representation
59         [~, ind] = max(Y); % one-hot encoding
60         ind = ind(1); % just in case there are multiple maxima
61
62         % keep the labels
63         trueLabels(letterNr, it) = letterNr;
64         simLabels(letterNr, it) = ind;
65
66         % check if it's the correct one
67         Nwrong(letterNr) = Nwrong(letterNr) + (ind ~= letterNr);
68     end
69 end
70 %icf(P) = sum(Nwrong) / (Nit*size(letters, 2)*size(letters, 1)); %
    normalize
71 icf(P) = sum(Nwrong) / (Nit*P); % normalize
72 fprintf('%i patterns give a normalized image error of %f\n', P, icf
    (P));
73 end
74
75 figure;
76 plot(1:P, icf, 'r-', 'LineWidth', 2);
77 xlabel('Number of patterns stored'); ylabel('Pixel error');
78 savefig('alternative_error.fig');
79
80 % THIS PART IS TOO INTENSIVE FOR MY PC
81 % %% CASE OF 10 PATTERNS

```

```

82 % letters = allLetters(:,1:10);
83 % numDist = 100; % number of distorted images per letter
84 % X = zeros(size(letters,1),size(letters,2)*numDist);
85 % T = zeros(size(letters,2),size(letters,2)*numDist);
86 % for il = 1:size(letters,2)
87 %     letter = letters(:,il);
88 %     for id = 1:numDist
89 %         X(:,(il-1)*numDist+id) = DistortImage(letter);
90 %         T(il,(il-1)*numDist+id) = 1; % only that one is 1
91 %     end
92 % end
93 %
94 % % create network with lots of neurons
95 % net = feedforwardnet(50);
96 % net.layers{1}.transferFcn = 'tansig';
97 % net.layers{2}.transferFcn = 'softmax';
98 % net = train(net, X, T, 'UseParallel', 'yes');
99 %
100 % Nit = 1000;
101 % % store number of wrong results
102 % Nwrong = zeros(1,size(letters,2));
103 % % we will make the confusion matrix later
104 % trueLabels = zeros(size(letters,2),Nit);
105 % simLabels = zeros(size(letters,2),Nit);
106 %
107 % % loop over Nit distorted images per letter and calculate the error
108 % parfor letterNr = 1:size(letters,2)
109 %     letter = letters(:,letterNr);
110 %     fprintf('Starting with letter %i out of %i...\n',letterNr,size(
111 %         letters,2))
112 %     for it = 1:Nit
113 %         distImage = DistortImage(letter);
114 %         % use the net to retrieve
115 %         [Y,~,~] = net(distImage);
116 %         % one-hot representation
117 %         [~,ind] = max(Y); % one-hot encoding
118 %         ind = ind(1); % just in case there are multiple maxima
119 %
120 %         % keep the labels
121 %         trueLabels(letterNr,it) = letterNr;
122 %         simLabels(letterNr,it) = ind;
123 %
124 %         % check if it's the correct one
125 %
126 %         Nwrong(letterNr) = Nwrong(letterNr) + (ind ~= letterNr);
127 %     end
128 % end
129 %
130 %
131 % fprintf('%i patterns give a normalized image error of %f\n',25,sum(

```

```
        Nwrong) / (Nit*25));  
132 %  
133 % % confusion matrix  
134 % trueLabels = reshape(trueLabels,[1 size(trueLabels,1)*size(  
        trueLabels,2)]);  
135 % simLabels = reshape(simLabels,[1 size(simLabels,1)*size(simLabels  
        ,2)]);  
136 % plotconfusion(ind2vec(trueLabels),ind2vec(simLabels))
```